

Lecture 1

Terraform basic blocks:

1. Terraform block
2. Provider block
3. Resource block

Terraform block: Terraform block contains: required versions, provider requirements and configure: some behaviour and terraform backend(Terraform state). This is the first step for creating anything with terraform.

Provider block: It relies on providers to interact with remote system(cloud), declare them(provider) to install and use them and lastly provider configuration belong to Root Module.

Resource Block: Each block describes one or more infra objects.

Resource syntax: How resource is declared

Resource Behavior: How Terraform handles resource declarations

Provisioners: Can configure resource post-creation actions

Lecture2

Terraform block

- Within a terraform block, only constant values can be used.
- Can't use any built in functions of terraform here also need to write value of the respective arguments, can't refer any variables.
- **Terraform block is introduced from 0.13 onwards**

Terraform 0.13 and later:

```
terraform {
  required_providers {
    aws = {
      source  = "hashicorp/aws"
      version = "~> 3.0"
    }
  }
}

# Configure the AWS Provider
provider "aws" {
  region = "us-east-1"
}

# Create a VPC
resource "aws_vpc" "example" {
  cidr_block = "10.0.0.0/16"
}
```

- Example:

- **Terraform block required:**

- required version. (need to match with local cli version else error, need to lock the version)

```
required_version = "~> 0.14.3"
```

- Required_providers: Will install during the **terraform init** command

```
backend "s3" {
  bucket = "mybucket"
  key    = "path/to/my/key"
  region = "us-east-1"
}
```

- Backend: Storing the state information of our configuration.

Storing the state information of our configuration.

- Terraform block optional:
- Experimental language features: This is experimental.

```
experiments = [ example ]
```

- Passing Metadata to Providers: More advanced terraform feature.

```
provider_meta "my-provider" {  
    hello = "world"  
}
```

etc.

- **Random Provider:** The "random" provider allows the use of randomness within Terraform configurations. This is a logical provider, which means that it works entirely within Terraform's logic, and doesn't interact with any other services.
 - Where random providers are used: Generate unique resource name, testing and simulation, load balancing and routing etc.

Lecture 3

Creation of terraform block.

Code:

1.

```
terraform {  
    required_version = "1.5.4"  
}
```

Won't work.

```
terraform {  
    required_version = "~> 1.5.4"  
}
```

Will work because it is allowing the right most version as I have already 1.5.5 so this syntax will work.

Because it ensure the latest minor version upgrade. It allows anything >1.5 but less than 1.6

Version Constraint Syntax

Terraform's syntax for version constraints is very similar to the syntax used by other dependency management systems.

version = ">= 1.2.0, < 2.0.0"

A version constraint is a [string literal](#) containing one or more conditions, which are separated by commas.

Each condition consists of an operator and a version number.

Version numbers should be a series of numbers separated by periods (like [1.2.0](#)), optionally with a suffix to indicate a beta release.

The following operators are valid:

- `=` (or no operator): Allows only one exact version number. Cannot be combined with other conditions.

- `!=`: Excludes an exact version number.
- `>=, <, <=`: Comparisons against a specified version, allowing versions for which the comparison is true. "Greater-than" requests newer versions, and "less-than" requests older versions.
- `~>`: Allows only the *rightmost* version component to increment. For example, to allow new patch releases within a specific minor release, use the full version number: `~> 1.0.4` will allow installation of `1.0.5` and `1.0.10` but not `1.1.0`. This is usually called the pessimistic constraint operator.(Important)

Lecture 4

It specifies all of the required providers by the current module, mention the source form which they need to download the provider and also ensures that provider constraints are there.

Latest aws provider version:

The screenshot shows a dropdown menu with 'Version 5.12.0' selected. Below it is a link 'Latest Version'. The main content area is titled 'LATEST VERSION' and lists several versions of Terraform:

- Version 5.12.0** (marked with a checkmark) - Published 6 days ago
- Version 5.11.0** - Published 13 days ago
- Version 5.10.0** - Published 20 days ago
- Version 5.9.0** - Published a month ago
- Version 5.8.0** - Published a month ago

[View all versions](#)

and we can see that the latest one is 5.12.0 and **terraform init** is installing this version.

Caution regarding required_version: always use the complete minor version. Because if we write anything like: “`~>0.14`” it will search for upgraded version for 14(15, 16..) and there is some major changes. So Always use complete minor version. Like: `0.14.0` then it will search for `0.14.1, 0.14.2....`

terraform init -upgrade will upgrade or downgrade based on the condition provided.

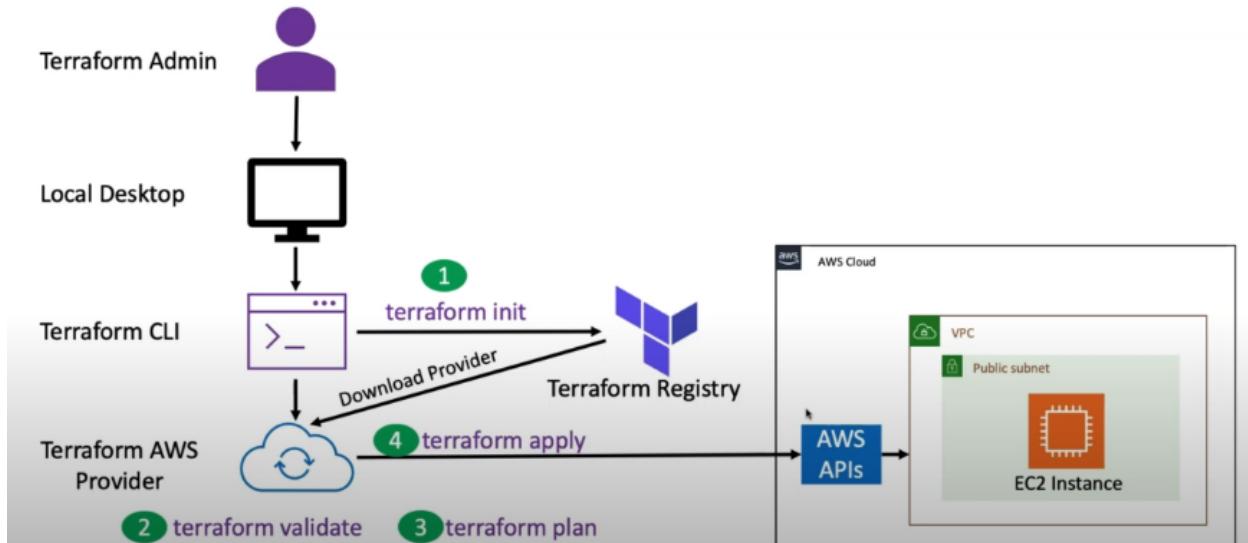
terraform version variation and check

```
/*
Play with Terraform Version
required_version = "~> 0.14.3"
required_version = "= 0.14.4"
required_version = ">= 0.13"
required_version = "= 0.13"
required_version = "~> 0.13"

Play with Provider Version
version = "~> 3.0"
version = ">= 3.0.0, < 3.1.0"
version = ">= 3.0.0, <= 3.1.0"
version = "~> 2.0"
version = ">= 3.0"
*/
```

Lecture 6

Terraform Providers



So the provider will communicate with the AWS APIs and those API's will provision the ec2 instance. The same as terraform apply is done during terraform destroy.

Without providers terraform can't manage any infrastructure. And these are distributed separately from the terraform and each provider has its own release cycle and version number.

Terraform got 3443 providers.

Terraform provider block:

Profile = "default"

It will use the default configuration from local cli or we can use access key, secret access key like static information.

Dependency lock file: Have the provider related information including version.

Local name for required providers:

```
# Terraform Block
terraform {
  required_version = "~> 0.14.3"
  required_providers {
    aws = {
      source = "hashicorp/aws"
      version = "~> 3.0"
    }
  }
}
```

Here aws is the local name and it is named after source="hashicorp/**aws**" It is recommended to use this name but any other local name can also be created.

Local names are module specific and should be unique per-module. And if I want to refer outside of the provider, then I can refer using the local

name.

Source: primary location from where we can download the provider.

It got 3 part:

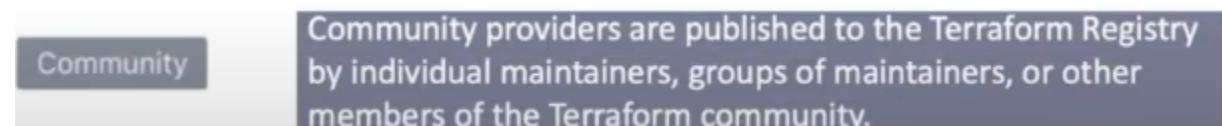
[<HOSTNAME>/]<NAMESPACE>/<TYPE>

Example: registry.terraform.io/hashicorp/aws

But registry name is optional when we are using terraform public registry.

Registry contains 4 things: Providers, Modules, Policy Libraries and Run Tasks.

Provider got badges(provided by terraform=Official, Verified, Community and archived(No longer maintained by hashicorp))



Lecture 7

Terraform needs access to respective cloud provider. There are several ways:

Configuration for the AWS Provider can be derived from several sources, which are applied in the following order:

1. Parameters in the provider configuration(Static credential is not recommended)
2. [Environment variables](#)(aws configure)
3. [Shared credentials files](#)
4. Shared configuration files
5. Container credentials

6. Instance profile credentials and region

Here first 3 are the main. **And already saved files are in .aws/credentials folder.**

Lecture 8

```
resource "aws_vpc" "main" {
  cidr_block = "10.0.0.0/16"(required)
}
```

Here aws_vpc is mandatory and main the convenient label to refer this vpc object inside terraform.

In terraform init = .terraform.lock.hcl file is created with .terraform folder which got provider version inside it.

Terraform tf state file: Terraform configuration database. Resources created in the aws cloud, it's equivalent information will be stored here. This file is created at terraform apply.

Lecture 9

Multiple provider: Define multiple configuration for the same provider and we can use on a per-resource basis.

Difference is to use alias.

Way to refer them: <PROVIDER NAME>.<ALIAS>

```
resource "aws_vpc" "my_vpc" {
  cidr_block      = "192.168.0.0/26"
  tags = {
    Name = "main"
  }
}
```

```

resource "aws_vpc" "my_vpc-us-east-2" {
  cidr_block      = "192.168.0.0/26"
  provider = aws.aws-east-2

  tags = {
    Name = "vpc-us-east-2"
  }
}

```

Here provider is a meta argument.

```

provider "aws" {
  profile = "default"
  region = "us-east-1"
}

provider "aws" {
  profile = "default"
  region = "us-east-2"
  alias = "aws-east-2"
}

```

Lecture 11

Terraform dependency lock file.

Terraform configuration refers to 2 different kinds of external dependency that come from outside of it's own codebase.

They are: **Providers and Modules**.

Dependency lock file works towards only provider version locking.
After selecting a specific version of each dependency using Version Constraints Terraform remembers the decision it made in a dependency lock file so that it can make the same decision again in future.

Inserted into github so that this version will be automatically downloaded in the users directory.

Lock file currently tracks only Provider dependencies.

Hashes in the lock file will be created during the **terraform init**.

If terraform didn't find any lock file it downloads the latest versions that fulfills the version constraints, but if not then terraform will use existing lock file and will install same provider version always.

Lecture 15

Terraform resource syntax introduction

Resource template:

```
# Template
<BLOCK TYPE> "<BLOCK LABEL>" "<BLOCK LABEL>" {
    # Block body
    <IDENTIFIER> = <EXPRESSION> # Argument
}
```

Top level blocks: resource, provider, variables etc.

Block inside block: tags etc.

Resources have 2 block labels. And variable have 1 block level.

Arguments: resource level details. They reside in the block body.

First block label determines the kind of infrastructure object it manages and what arguments and other attributes the resource supports = **Resource Type**

Second block label is used to refer this resource within module. But it got no scope outside this module = **Resource Local Name**

Local name needs to be unique also.

Can be used with any resource to chance the behaviour of resources = **Meta-Argument**. Othets are argument.

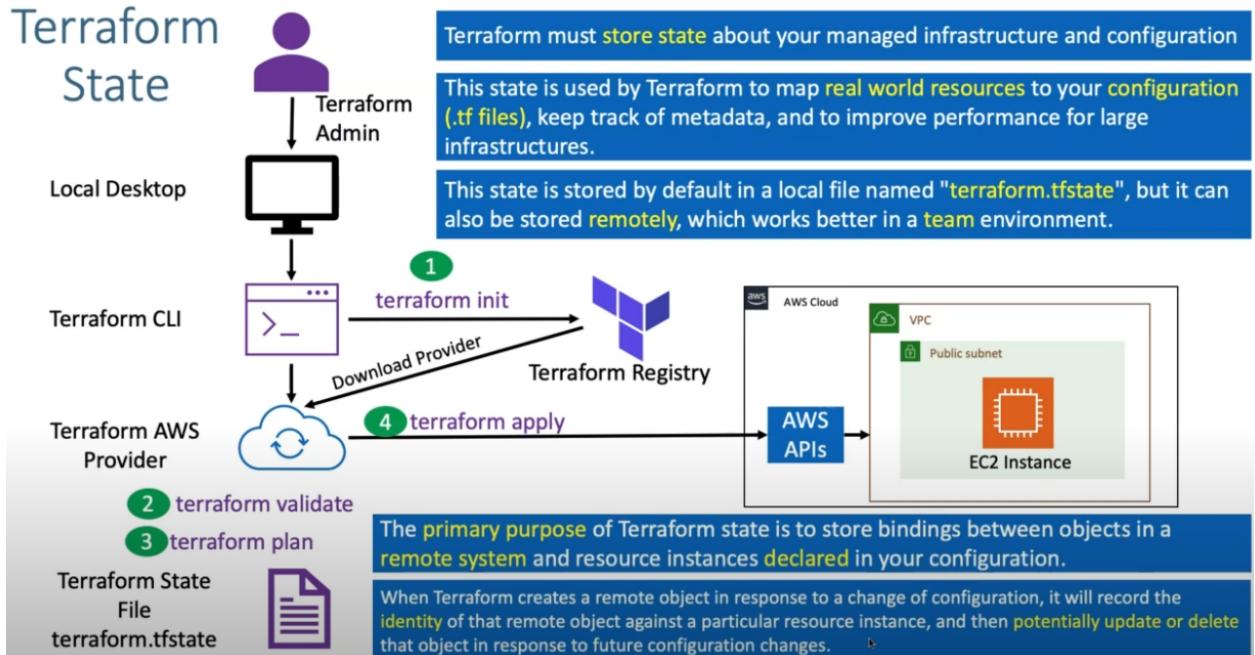
Lecture 16

Resource Behavior



Lecture 18

Terraform State



Terraform used the selected providers to generate:
~ update in-place

Terraform will perform the following actions:

- Desired State: Local Terraform Manifest (All `*.tf` files)
- Current State: Real Resources present in your cloud

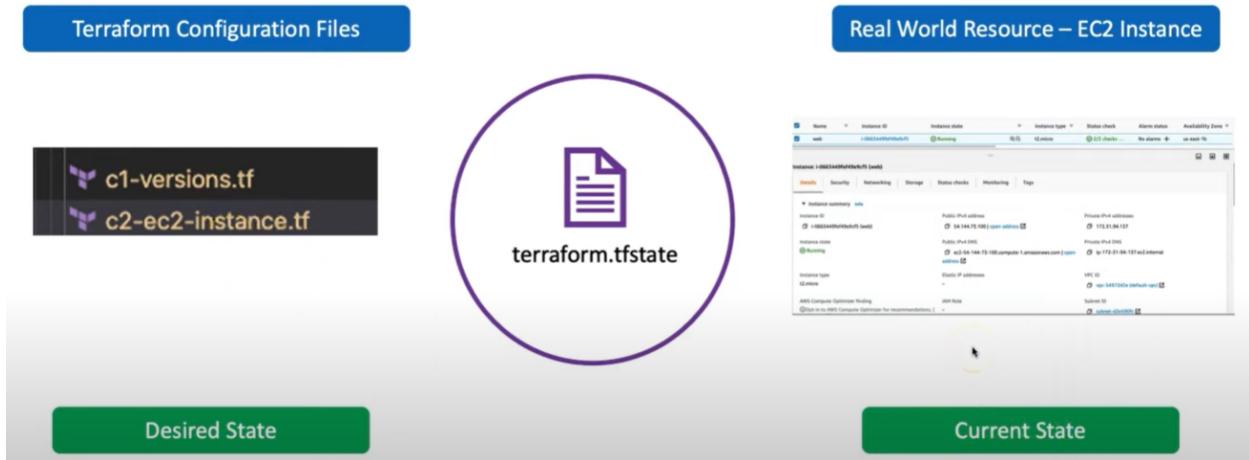
Lecture 20:

Destroy and recreate resources

```
# aws_instance.my-ec2-vm must be replaced
-/+ resource "aws_instance" "my-ec2-vm" {
```

Because in case of available zone change, it will first destroy and then create into another zone. So, destroy and create.

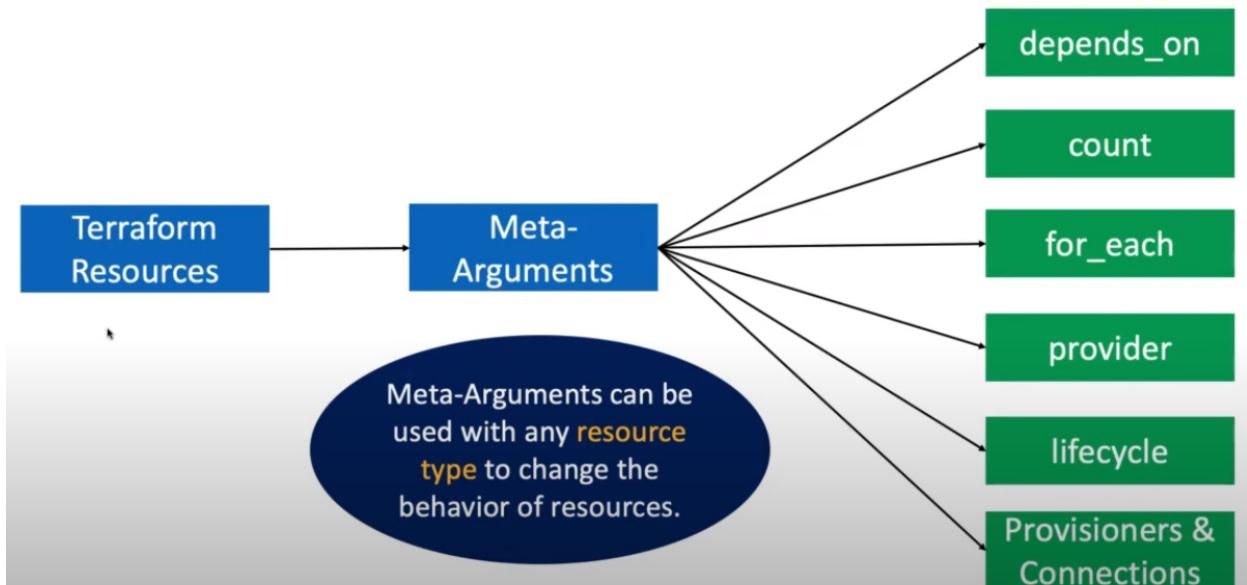
Lecture 21



Lecture 22

Meta argument.

Resource Meta-Arguments



Resource Meta-Arguments

| | |
|----------------------------|---|
| depends_on | To handle hidden resource or module dependencies that Terraform can't automatically infer. |
| count | For creating multiple resource instances according to a count |
| for_each | To create multiple instances according to a map , or set of strings |
| provider | For selecting a non-default provider configuration |
| lifecycle | Standard Resource behavior can be altered using special nested lifecycle block within a resource block body |
| Provisioners & Connections | For taking extra actions after resource creation (Example: install some app on server or do something on local desktop after resource is created at remote destination) |



Lecture 29

```
#user_data = file("apache-install.sh")
```

Lecture 30

Count Meta Argument.

The `count` Object

In blocks where `count` is set, an additional `count` object is available in expressions, so you can modify the configuration of each instance. This object has one attribute:

- `count.index` — The distinct index number (starting with `0`) corresponding to this instance.

Using Expressions in `count`

The `count` meta-argument accepts numeric [expressions](#). However, unlike most arguments, the `count` value must be known *before* Terraform performs any remote resource actions. This means `count` can't refer to any resource attributes that aren't known until after a configuration is applied (such as a unique ID generated by the remote API when an object is created).

When to Use `for_each` Instead of `count`

If your instances are almost identical, `count` is appropriate. If some of their arguments need distinct values that can't be directly derived from an integer, it's safer to use `for_each`.

Resource Meta-Arguments – count

If a **resource or module** block includes a **count** argument whose value is a **whole number**, Terraform will create that **many instances**.

count.index: The distinct index number (starting with 0) corresponding to this instance.

Each instance has a **distinct infrastructure object associated with it**, and each is separately **created, updated, or destroyed** when the configuration is applied.

Resource
Meta-Argument
count

When count is set, Terraform **distinguishes** between the block itself and the multiple resource or module instances associated with it. Instances are identified by an **index number, starting with 0**. `aws_instance.myvm[0]`

The count meta-argument accepts **numeric expressions**. The count value must be known **before** Terraform performs any remote resource actions.

Module support for count was added in **Terraform 0.13**, and previous versions can only use it with **resources**.

A given resource or module block **cannot** use both **count** and **for_each**

Lecture 32

Resource Meta-Arguments – for_each

If a **resource or module** block includes a **for_each** argument whose value is a **map** or a **set of strings**, Terraform will create **one instance for each member** of that map or set.

A given resource or module block **cannot** use both **count** and **for_each**

For set of Strings, `each.key = each.value`
`for_each = toset(["Jack", "James"])`
`each.key = Jack`
`each.key = James`

Each instance has a **distinct infrastructure object associated with it**, and each is separately **created, updated, or destroyed** when the configuration is applied.

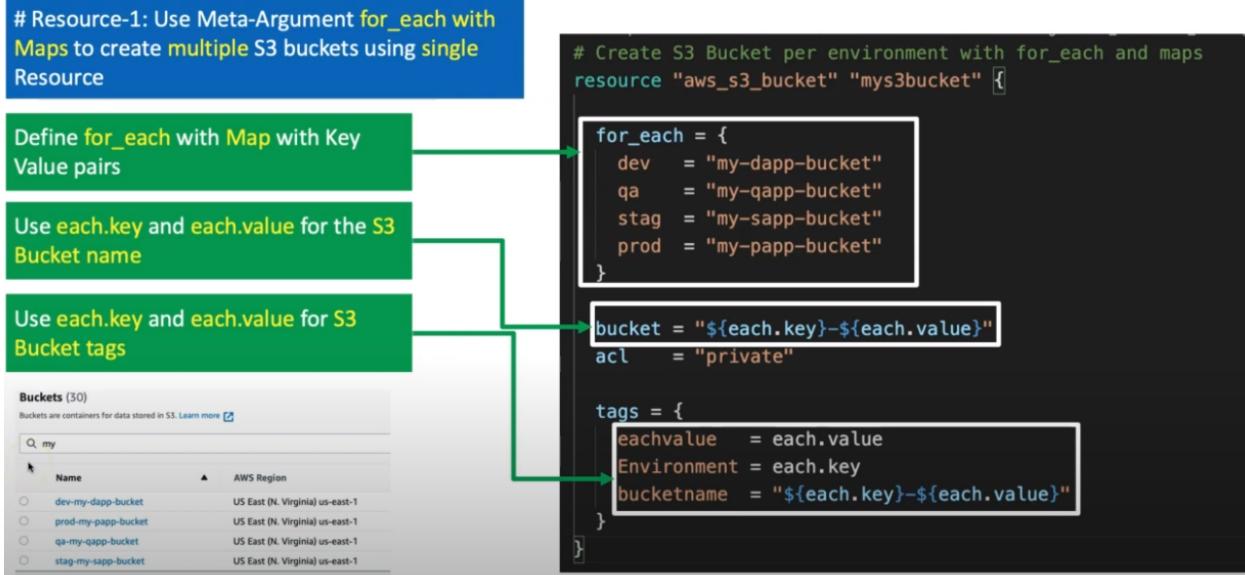
Resource
Meta-Argument
for_each

For Maps, we use `each.key & each.value`
`for_each = {`
 `dev = "my-dapp-bucket"`
`}`
`each.key = dev`
`each.value = my-dapp-bucket`

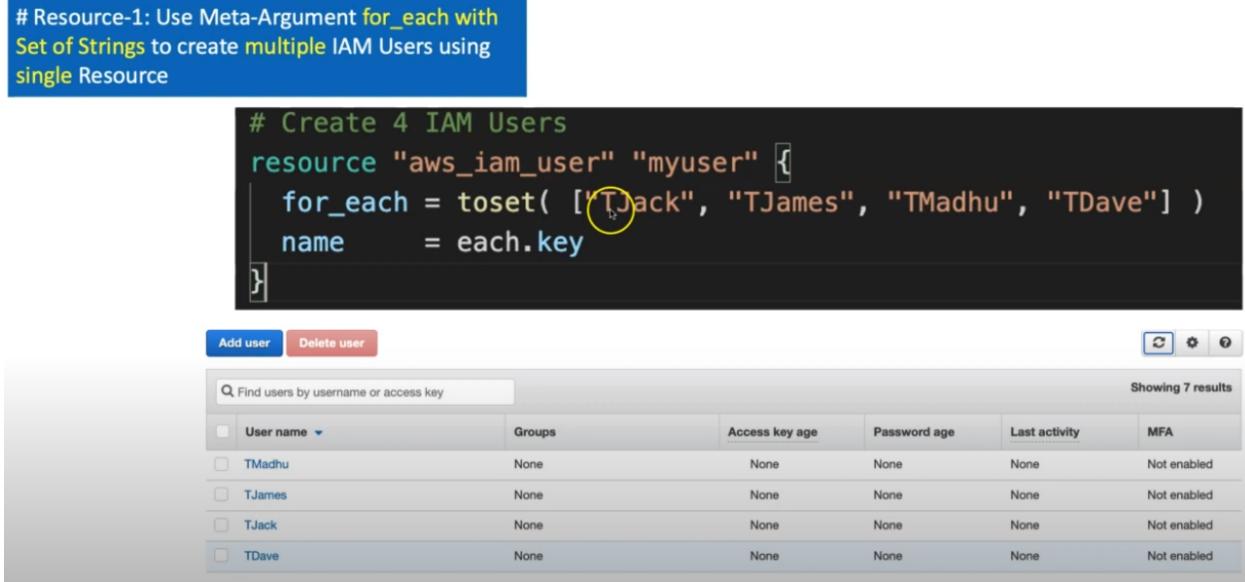
In blocks where **for_each** is set, an additional **each** object is available in expressions, so you can modify the configuration of each instance.
each.key — The map key (or set member) corresponding to this instance.
each.value — The map value corresponding to this instance. (If a set was provided, this is the same as **each.key**.)

Module support for **for_each** was added in **Terraform 0.13**, and previous versions can only use it with **resources**.

Usecase-1: for_each Maps



Usecase-2: for_each Set of Strings (toset)



Lecture 33

```
resource "aws_s3_bucket" "tf_s3" {

for_each = {

dev    = "my-dapp-bucket"

qa     = "my-qapp-bucket"

stag   = "my-sapp-bucket"

prod   = "my-papp-bucket"

}

#  each.key

#  each.value

bucket = each.value

#  acl = "private"

tags = {

# Name          = "My bucket"

Environment = each.key

bucketname  = "${each.key}-${each.value}"

eachvalue   = each.value

}

}
```

```
}
```

```
# resource "aws_s3_bucket_acl" "name" {
```

```
# }
```

```
# resource "aws_instance" "tf_ec2" {
```

```
#   ami                      = "ami-08a52ddb321b32a8c"
```

```
#   instance_type            = "t2.micro"
```

```
#   key_name                = "tf-key"
```

```
#   subnet_id               = aws_subnet.pubsub_1.id
```

```
#   vpc_security_group_ids = [aws_security_group.pubsub-1-sg.id]
```

```
#   availability_zone       = "us-east-1a"
```

```
#   tags = {
```

```
#     Name = "E2-pubsub-1"
```

```
#     tag1 = "New tag"
```

```
#   }
```



```
#   # here EOF is used but we can also use file function.
```

```
#   user_data = <<-EOF
```

```
#       #!/bin/bash
```

```
#       sudo yum update -y
```

```
#       sudo yum install -y httpd
```

```
#       sudo systemctl start httpd
```

```
#           sudo systemctl enable httpd

#           echo "<h1>Welcome : Output form EOF</h1>" | sudo tee
/var/www/html/index.html

#           EOF

# }

resource "aws_instance" "tf_ec2" {

  ami          = "ami-08a52ddb321b32a8c" #Amazon Linux of us-east-1
region

  instance_type = "t2.micro"

  # key_name          = "tf-key"

  subnet_id = aws_subnet.pubsub_1.id

  # count      = 5

  # vpc_security_group_ids = [aws_security_group.pubsub-1-sg.id]

  # availability_zone     = "us-east-1a"

  tags = {

    Name = "E2-pubsub-1"

    # Name = "E2-pubsub-1-${count.index}"

    # tag1 = "New tag"

  }

}
```

```
---  
  
terraform {  
  
  # required_version = "~> 1.5"  
  
  required_version = "~> 0"  
  
  required_providers {  
  
    aws = {  
  
      source  = "hashicorp/aws"  
  
      version = "= 5.12.0"  
  
    }  
  
    random = {  
  
      source  = "hashicorp/random"  
  
      version = "3.4.3"  
  
    }  
  
  }  
  
}  
  
  
# provider "aws" {  
  
#   profile = "default"  
  
#   region = "us-ease-1"  
  
# }
```

```
provider "aws" {

  profile = "default"

  region  = "us-east-1"

}

# provider "aws" {

#   profile = "default"

#   region  = "us-east-2"

#   alias   = "aws-east-2"

# }
```

```
resource "aws_eip" "pubsub-1-eip" {

  instance = aws_instance.tf_ec2.id

  domain   = "vpc"

  #Meta argument

#  depends_on = [aws_internet_gateway.gw]

}
```

```
resource "aws_vpc" "my_vpc" {

cidr_block = "192.168.0.0/26"

tags = {

Name = "main"

}

}

# resource "aws_vpc" "my_vpc-us-east-2" {

#   cidr_block      = "192.168.0.0/26"

#   provider = aws.aws-east-2

# }

#   tags = {

#     Name = "vpc-us-east-2"

#   }

# }
```

#Creating subnet

```
resource "aws_subnet" "pubsub_1" {

vpc_id          = aws_vpc.my_vpc.id
cidr_block      = "192.168.0.0/27" #first public subnet
availability_zone = "us-east-1a"
```

```
map_public_ip_on_launch = true #ensures the public ip of an instance

tags = {

    Name = "Public Subnet 1"

}

}

# Creating internet gateway

resource "aws_internet_gateway" "gw" {

vpc_id = aws_vpc.my_vpc.id

tags = {

    Name = "pubsub-1-igw"

}

}

# Creating route table

resource "aws_route_table" "pubsub-1-rt" {

vpc_id = aws_vpc.my_vpc.id #required

route {

cidr_block = "0.0.0.0/0"

gateway_id = aws_internet_gateway.gw.id
}
```

```
}

tags = {

    Name = "pubsub-1-rt"
}

}

# Subnet association to route table

resource "aws_route_table_association" "pubsub_1_association" {

    subnet_id      = aws_subnet.pubsub_1.id

    route_table_id = aws_route_table.pubsub-1-rt.id
}

# Creating security group

resource "aws_security_group" "pubsub-1-sg" {

    name = "pubsub-1-sg"

    # description = "Allow TLS inbound traffic"

    vpc_id = aws_vpc.my_vpc.id

    ingress {

        description = "Allow port 22"

        from_port   = 22

        to_port     = 22
    }
}
```

```
protocol      = "tcp"

cidr_blocks = ["0.0.0.0/0"]

}

ingress {

description = "Allow port 80"

from_port   = 80

to_port     = 80

protocol    = "tcp"

cidr_blocks = ["0.0.0.0/0"]

}

egress {

from_port   = 0

to_port     = 0

protocol    = "-1"

cidr_blocks = ["0.0.0.0/0"]

# ipv6_cidr_blocks = [":/:0"]

}

tags = {

Name = "Allow port 22 and 80"

}
```

```
}
```

```
resource "aws_instance" "tf_ec2" {
  ami      = "ami-0ccabb5f82d4c9af5"
  instance_type = "t2.micro"
  availability_zone = "us-east-2a"
  tags = {
    Name = "terraform_ec2"
    tag1 = "New tag"
  }
}
```

```
resource "aws_vpc" "my_vpc" {
  cidr_block      = "192.168.0.0/26"

  tags = {
    Name = "main"
  }
}

# resource "aws_vpc" "my_vpc-us-east-2" {
#   cidr_block      = "192.168.0.0/26"
#   provider = aws.aws-east-2

#   tags = {
```

```
#   Name = "vpc-us-east-2"
#
# }
```

Lecture 34

Toset example:

```
resource "aws_iam_role" "toset_FOREACH" {
  for_each = toset(["TJack", "Tjames", "Tmahhdu"])
  name      = each.key
}
```

Lecture 33

lifecycle is a **nested block** that can appear within a resource block

Resource
Meta-Argument
lifecycle

The lifecycle block and its contents are **meta-arguments**, available for all resource blocks regardless of type

create_before_destroy

prevent_destroy

ignore_changes

```
Create EC2 Instance
resource "aws_instance" "web" [
  ami = "ami-0915bcb5fa77e4892" # A
  instance_type = "t2.micro"
  availability_zone = "us-east-1a"
  #availability_zone = "us-east-1b"
  tags = {
    "Name" = "web-1"
  }
  lifecycle {
    create_before_destroy = true
  }
]
```

```
# Create EC2 Instance
resource "aws_instance" "web" [
  ami = "ami-0915bcb5fa77e4892"
  instance_type = "t2.micro"
  tags = {
    "Name" = "web-2"
  }
  lifecycle {
    prevent_destroy = true # Def
  }
]
```

```
# Create EC2 Instance
resource "aws_instance" "web" [
  ami = "ami-0915bcb5fa77e4892"
  instance_type = "t2.micro"
  tags = {
    "Name" = "web-3"
  }
  lifecycle [
    ignore_changes = [
      # Ignore changes to tags
      # updates these based on
      tags,
    ]
  ]
]
```

Lecture 38

lifecycle is a **nested block** that can appear within a resource block

Resource Meta-Argument lifecycle

The lifecycle block and its contents are **meta-arguments**, available for **all** resource blocks regardless of **type**.

create_before_destroy

```
# Create EC2 Instance
resource "aws_instance" "web" {
  ami = "ami-0915bcb5fa77e4892" # Amazon Linux
  instance_type = "t2.micro"
  availability_zone = "us-east-1a"
  #availability_zone = "us-east-1b"
  tags = {
    "Name" = "web-1"
  }
  lifecycle {
    create_before_destroy = true
  }
}
```

prevent_destroy

```
# Create EC2 Instance
resource "aws_instance" "web" {
  ami = "ami-0915bcb5fa77e4892"
  instance_type = "t2.micro"
  tags = {
    "Name" = "web-2"
  }
  lifecycle {
    prevent_destroy = true # Default is false
  }
}
```

ignore_changes

```
# Create EC2 Instance
resource "aws_instance" "web" {
  ami = "ami-0915bcb5fa77e4892"
  instance_type = "t2.micro"
  tags = {
    "Name" = "web-3"
  }
  lifecycle {
    ignore_changes = [
      # Ignore changes to tags,
      # updates these based on tags,
    ]
  }
}
```

Apply complete! Resources: 1 added, 0 changed, 0 destroyed.

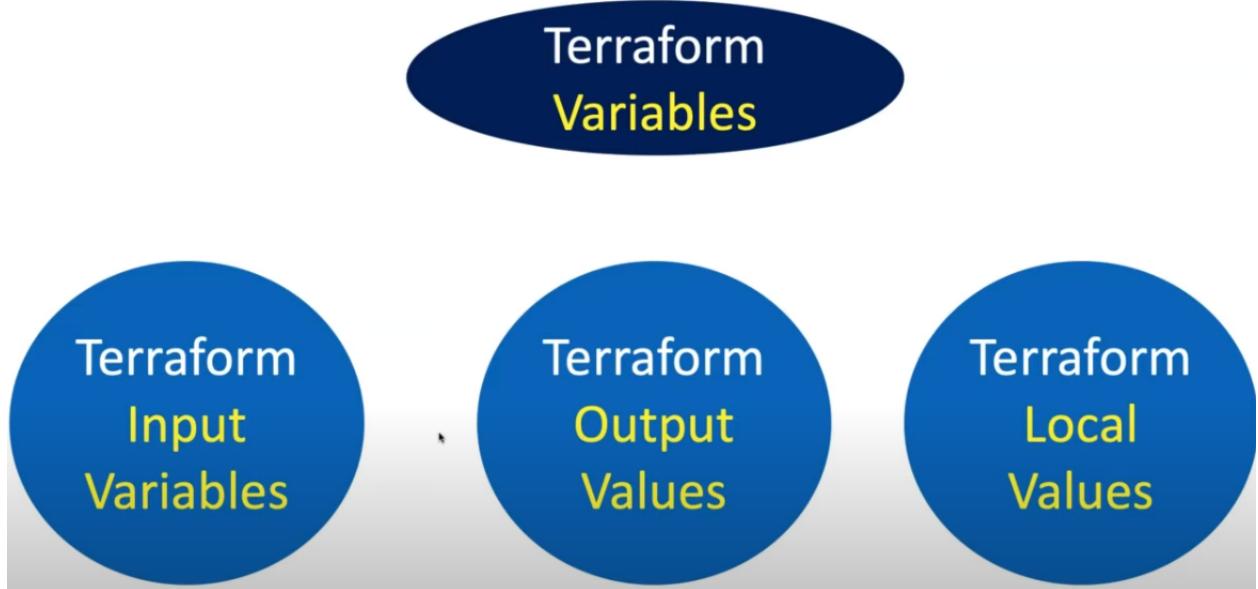
Kalyans-MacBook-Pro:v2-prevent_destroy kdaida\$ terraform destroy -auto-approve

Error: Instance cannot be destroyed

on c2-ec2-instance.tf line 2:
2: resource "aws_instance" "web" {

```
1 # Create EC2 Instance
2 resource "aws_instance" "web" {
3   ami = "ami-0915bcb5fa77e4892" # Amazon Linux
4   instance_type = "t2.micro"
5   tags = {
6     "Name" = "web-2"
7   }
8   /*
9   lifecycle {
10     prevent_destroy = true # Default is false
11   }
12 }
13
14
```

Lecture 41



Why S3: S3 buckets aim to help enterprises and individuals achieve their data backup and delivery needs. S3 buckets also enable a large amount of data to be stored and accessed later through cloud storage. Most data stored by enterprises in S3 buckets is for big data analytics, disaster recovery, dynamic websites, and user-generated content, among other uses. Some enterprises also use S3 buckets to host static HTML websites and dynamically complex web applications.

`terraform destroy` does not delete the S3 Bucket ACL but does remove the resource from Terraform state.

What is an S3 bucket ACL?

Amazon S3 access control lists (ACLs) enable you to manage access to S3 buckets and objects. Every S3 bucket and object has an ACL attached to it as a subresource. The ACLs define which AWS accounts or groups are granted access along with the type of access.

- `bucket` - (Required) Name of the bucket that you want to associate this access point with.
- `rule` - (Required) Configuration block(s) with Ownership Controls rules. Detailed below.

rule Configuration Block

The following arguments are required:

- `object_ownership` - (Required) Object ownership. Valid values:
`BucketOwnerPreferred`, `ObjectWriter` or `BucketOwnerEnforced`
 - `BucketOwnerPreferred` - Objects uploaded to the bucket change ownership to the bucket owner if the objects are uploaded with the `bucket-owner-full-control` canned ACL.
 - `ObjectWriter` - Uploading account will own the object if the object is uploaded with the `bucket-owner-full-control` canned ACL.
 - `BucketOwnerEnforced` - Bucket owner automatically owns and has full control over every object in the bucket. ACLs no longer affect permissions to data in the S3 bucket.

Why bucket ownership control block is used:

The `aws_s3_bucket_ownership_controls` resource is used to manage the ownership controls of an S3 bucket. Ownership controls are a feature of S3 that allow you to control who owns objects that are uploaded to the bucket.

By default, an object that is uploaded to an S3 bucket is owned by the AWS account that uploaded the object, not the owner of the bucket. This can lead to situations where the bucket

owner does not have access to objects in their own bucket, or where they are not able to manage access permissions for those objects.

The `aws_s3_bucket_ownership_controls` resource allows you to change this behavior. By setting the `object_ownership` attribute to `BucketOwnerPreferred` or `BucketOwnerEnforced`, you can ensure that the bucket owner will own all new objects that are uploaded to the bucket, regardless of who uploaded them.

This is particularly useful in situations where you want to ensure that the bucket owner has full control over all objects in the bucket, or where you want to centralize the management of access permissions.

Why ownership and acl is used:(Overall)

ACLs (Access Control Lists) and ownership controls are two mechanisms provided by AWS S3 to manage access permissions and ownership of objects within S3 buckets.

1. ACLs: ACLs allow you to control access to your S3 bucket and objects at the individual object level. With ACLs, you can grant or deny read and write permissions to specific AWS accounts or groups. ACLs provide a more granular level of access control compared to bucket policies or IAM policies, as they can be applied to individual objects within a bucket.

Key points about ACLs:

- ACLs define who can access the bucket or object and the level of access they have, such as read or write.
 - ACLs are a legacy access control mechanism in S3.
 - By default, ACLs are disabled, and the bucket owner automatically has full control over every object in the bucket.
 - However, ACLs can be useful in scenarios where you need specific access control for individual objects within the bucket.
2. Ownership Controls: Ownership controls help ensure that the bucket owner has full control over the objects in their S3 bucket. By default, the AWS account that uploads an object owns that object. This can lead to situations where the bucket owner does not have access to their own objects or cannot manage access permissions for those objects.

Ownership controls address this issue by allowing the bucket owner to become the owner of every object uploaded to the bucket, regardless of who performed the upload. The most common ownership control rule is the "BucketOwnerPreferred" rule. When this rule is enabled, the bucket owner becomes the object owner, giving them complete control over the objects' access and permissions.

Key points about ownership controls:

- Ownership controls ensure that the bucket owner has full control over every object in the bucket.
- The "BucketOwnerPreferred" rule is commonly used to specify that the bucket owner should be the owner of every object uploaded to the bucket.
- Ownership controls provide a way to centralize ownership and access control within the bucket owner's AWS account.

```

Error: Insufficient rule blocks

on s3_bucket.tf line 26, in resource "aws_s3_bucket_ownership_controls" "example":
26: resource "aws_s3_bucket_ownership_controls" "example" {

At least 1 "rule" blocks are required.

$ terraform plan

Error: Missing required argument

on s3_bucket.tf line 40, in resource "aws_s3_bucket_acl" "tf_s3_acl":
40: resource "aws_s3_bucket_acl" "tf_s3_acl" {

The argument "bucket" is required, but no definition was found.

```

This problem occurs if

```
depends_on = [aws_s3_bucket_ownership_controls.example]
is not provided.
```

```

Error: creating S3 bucket ACL for prod-my-papp-bucket-keya: AccessControlListNotSupported: The bucket does not allow ACLs
status code: 400, request id: HE4CYQABJW63989B, host id: UA855t+LA+3Sr8TRRRj3VDEKQwP4Dd4Hdm7xx0gk7e3juFesj251chp2ui4T31HpsSl2cRoorws=
with aws_s3_bucket_acl.tf_s3_acl["prod"],
on s3_bucket.tf line 40, in resource "aws_s3_bucket_acl" "tf_s3_acl":
40: resource "aws_s3_bucket_acl" "tf_s3_acl" {

```

Lecture 42

Terraform Input Variables

Input variables serve as **parameters** for a Terraform module, allowing aspects of the module to be **customized** without altering the module's own source code, and allowing modules to be **shared** between different configurations.

Lecture 43

Setting the alias: `terraform`

Follow below steps:

1. Open the file `.bashrc` which is found in location
`C:\Users\USERNAME\.bashrc`
 If file `.bashrc` not exist then create it using below steps:

1. Open Command Prompt and goto `C:\Users\USERNAME\`.

2. Type command **notepad ~/ .bashrc**
It generates the **.bashrc** file.
 2. Add below sample commands of WP CLI, Git, Grunt & PHPCS etc.
-

```
# -----
# Git Command Aliases
# -----
alias ga='git add'
alias gaa='git add .'
alias gaaa='git add --all'

# -----
# WP CLI
# -----
alias wpthl='wp theme list'
alias wppll='wp plugin list'
```

Var region:

```
provider "aws" {
  profile = "default"
  # region  = "us-east-1"
  region = var.aws_region
}

variable "aws_region" {
  type = string
  default = "us-east-1"
}
```

```
egress {
  from_port    = 0
  to_port      = 0
  protocol     = "-1"
  cidr_blocks = ["0.0.0.0/0"]
  # ipv6_cidr_blocks = ["::/0"]
}
```

In the code snippet you provided, the egress rule allows all outbound traffic (from_port = 0, to_port = 0) with any protocol (protocol = -1) to any destination (cidr_blocks = ["0.0.0.0/0"]). This effectively allows all traffic to be sent from the associated instances to any destination IP address.

Lecture 45:

Input variables assigned when prompted

```
instance_type = var.ec2_instance_type #prompt example
```

```
$ tf plan
var.ec2_instance_type
  EC2 instance using variable

  Enter a value: t2.micro
```

```
instance_type = "t2.micro" -> null
```

Lecture 46

-var in CLI

```
BJIT@11728-Mahmuda-Akter-Keya MINGW64 /d/terraform tutorial
$ tf plan -var="ec2_instance_count=2" -var="ec2_instance_type=t2.micro"
```

```
tf plan -var="ec2_instance_count=2"
-var="ec2_instance_type=t2.micro"
```

```
BJIT@11728-Mahmuda-Akter-Keya MINGW64 /d/terraform tutorial
$ tf plan -var="ec2_instance_count=2" -var="ec2_instance_type=t2.micro" -out
v3out.plan
```

- **-out v3out.plan:** This flag is used to specify the output file name for the execution plan. In this case, the plan will be saved to a file named v3out.plan. -out=FILE option to save the generated plan to a file on disk.

```
BJIT@11728-Mahmuda-Akter-Keya MINGW64 /d/terraform tutorial
$ tf apply v3out.plan
```

- [REDACTED] to run the code

Lecture 47

Input variable with environment variable.

```
BJIT@11728-Mahmuda-Akter-Keya MINGW64 /d/terraform tutorial
$ export TF_VAR_ec2_instance_count=1
```

```
BJIT@11728-Mahmuda-Akter-Keya MINGW64 /d/terraform tutorial
$ export TF_VAR_ec2_instance_type=t2.micro
```

```
BJIT@11728-Mahmuda-Akter-Keya MINGW64 /d/terraform tutorial
$ echo $TF_VAR_ec2_instance_count, $TF_VAR_ec2_instance_type
1, t2.micro
```

Unsetting the environment variable:

```
BJIT@11728-Mahmuda-Akter-Keya MINGW64 /d/terraform tutorial
$ unset TF_VAR_ec2_instance_count
```

```
BJIT@11728-Mahmuda-Akter-Keya MINGW64 /d/terraform tutorial
$ unset TF_VAR_ec2_instance_type
```

```
BJIT@11728-Mahmuda-Akter-Keya MINGW64 /d/terraform tutorial
$ echo $TF_VAR_ec2_instance_count, $TF_VAR_ec2_instance_type
,
```

Lecture 48

Input from `terraform.tfvars` file.

- If the file name is `terraform.tfvars`, terraform will auto-load the variables present in this file by overriding the `default` values in `variables.tf`

```
Y  terraform.tfvars > abc ec2_instance_type
1   ec2_instance_count = 2
2   ec2_instance_type = "t3.micro"
```

Lecture 49

Input Variables Assign with `var` file argument. We use this for using separate `.tfvar` files.

Example

```
terraform plan -var-file="web.tfvars"
```

```
BJIT@11728-Mahmuda-Akter-Keya MINGW64 /d/terraform tutorial
$ tf plan -var-file="app.tfvars"
```

```
~ resource "aws_instance" "tf_ec2" {
    id                               = "i-0f1f5332ee4f30c8e"
    ~ instance_type                  = "t2.micro" -> "t3.micro"
}
```

```
BJIT@11728-Mahmuda-Akter-Keya MINGW64 /d/terraform tutorial
$ tf plan -var-file="web.tfvars"
```

```
~ instance_type                  = "t2.micro" -> "t3.small"
```

Lecture 50

Input variable with `.auto.tfvars`

```
Y  web.auto.tfvars > abc ec2_instance_type
1   ec2_instance_type = "t3.micro"
```

Whatever stays in this file will auto load during terraform plan and apply. Apply terraform plan

Output:

```
# aws_instance.tf_ec2[1] will be updated in-place
~ resource "aws_instance" "tf_ec2" {
    id                                = "i-0f1f5332ee4f30c8e"
    ~ instance_type                     = "t2.micro" -> "t3.micro"
    tags                               = {
        "Name" = "E2-pubsub-1"
    }
}
```

Lecture 51

- list (or tuple): a sequence of values, like ["us-west-1a", "us-west-1c"].
Elements in a list or tuple are identified by consecutive whole numbers, starting with zero.

```
#prompt example
variable "ec2_instance_type" {
    description = "EC2 instance using variable"
    # type      = string
    type = list(string)
    default = [ "t3.micro", "t3.small", "t3.large" ]
    # default = "t3.micro" # -var cli
}
```

and output

```
# aws_instance.tf_ec2[0] will be updated in-place
~ resource "aws_instance" "tf_ec2" {
    id                                = "i-0b718b8a564c62354"
    ~ instance_type                     = "t2.micro" -> "t3.micro"
    tags                               = {
        "Name" = "E2-pubsub-1"
    }
}
```

Lecture 52

Input Variables Complex Constructor of Type Map

```
instance_type = var.ec2_instance_type_map["small-apps"]
```

```
variable "ec2_instance_type_map" {
  description = "EC2 instance type"
  type = map(string)
  default = {
    "small-apps" = "t3.micro"
    "medium-apps" = "t3.medium"
    "big-apps" = "t3.large"
  }
}
```

Lecture 53

Custom validation rules in Variables

Terraform console:

The Terraform console is an interpreter that you can use to evaluate Terraform expressions and explore your Terraform project's state.

Scripting

```
$ echo "1+5" | terraform console
6
```

Terraform console command prompt

```
$ terraform console
> 1+5
6
```

Length function

- Length of a string

```
> length("😊😊😊⚡😊")
5
```

- Length of a list

```
> length(["a", "b", "c"])
3
```

- Length of Map

```
> length({"key" = "value"})
1
> length({"key1" = "value", "key2" = "value2"})
2
```

substr Function

substr(string, offset, length)

`substr` extracts a substring from a given string by offset and (maximum) length.

The offset index may be negative, in which case it is relative to the end of the given string. The length may be -1, in which case the remainder of the string after the given offset will be returned.

```
> substr("hello world", -5, -1)
```

world

If the length is greater than the length of the string, the substring will be the length of all remaining characters.

```
> substr("hello world", 6, 10)
```

world

Some example:

```
> substr("stack simplify", 1, 4)
"tack"
> substr("stack simplify", 0, 6)
"stack "
> substr("stack simplify", 0, 1)
"s"
> substr("stack simplify", 0,0)
...
> substr("stack simplify", 0, 10)
"stack simp"
```

Lecture 54

Custom validation rule

```
variable "ec2_ami_id" {
  type    = string
  default = "ami-08a52ddb321b32a8c"
  validation {
    condition = length(var.ec2_ami_id) > 4 && substr(var.ec2_ami_id, 0, 4) == "ami-"
    error_message = "The image_id value must be a valid AMI id, starting with \"ami-\"."
  }
}
```

Lecture 55

Protect sensitive input variable

Often we need to configure our infrastructure using sensitive or secret information such as usernames, passwords, API tokens, or Personally Identifiable Information (PII). When we do so, we need to ensure that you do not accidentally expose this data in CLI output, log output, or source control. Terraform provides several features to help avoid accidentally exposing sensitive data. In this example, we are working with rds database username and password as sensitive information.

Example:

```
variable "db_username" {
  description = "AWS RDS Database Administrator Username"
  type = string
  sensitive = true
}

variable "db_password" [
  description = "AWS RDS Database Administrator Password"
  type = string
  sensitive = true
]
```

```
rds-db.tf > ↗ resource "aws_db_instance" "db1"
1   resource "aws_db_instance" "db1" {
2     allocated_storage      = 5
3     engine                 = "mysql"
4     instance_class         = "db.t2.micro"
5     #   name                  = "mydb1"
6     engine_version         = "14.1"
7     username               = var.db_username
8     password               = var.db_password
9     skip_final_snapshot    = true
10 }
```

Lecture 56

Variable definition precedence

The above mechanisms for setting variables can be used together in any combination. If the same variable is assigned multiple values, Terraform uses the *last* value it finds, overriding any previous values. Note that the same variable cannot be assigned multiple values within a single source.

- Environment variables
- The `terraform.tfvars` file, if present.

- The `terraform.tfvars.json` file, if present.
- Any `*.auto.tfvars` or `*.auto.tfvars.json` files, processed in lexical order of their filenames.
- Any `-var` and `-var-file` options on the command line, in the order they are provided. (This includes variables set by a Terraform Cloud workspace.)

Lecture 57

Terraform file function

`file` reads the contents of a file at the given path and returns them as a string.

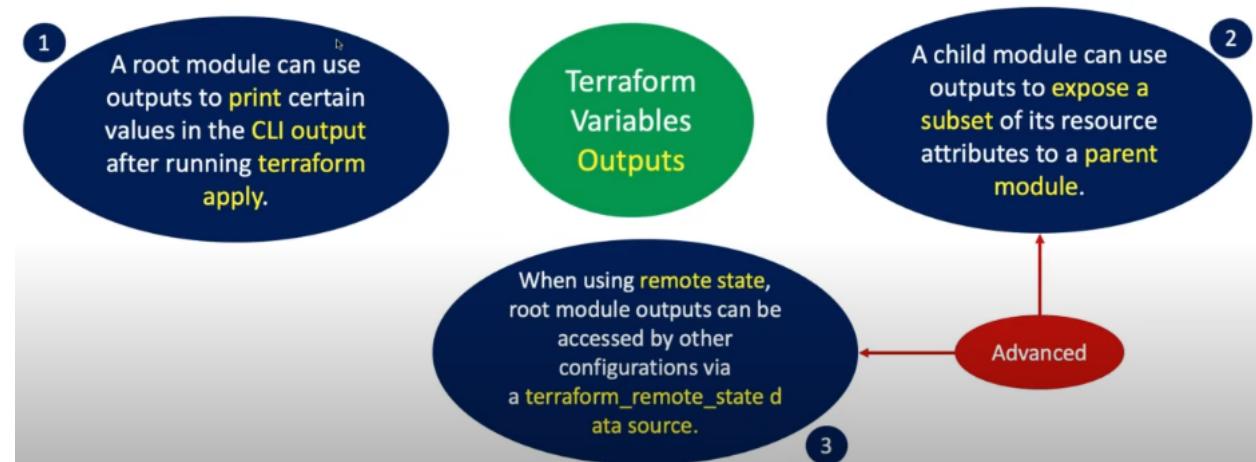
```
user_data = file("apache-install.sh")
```

Lecture 58

Terraform output values

Terraform Variables – Output Values

Output values are like the **return values** of a Terraform module and have several uses



Lecture 59

Attribute reference

Argument reference

Output

```
Apply complete! Resources: 8 added, 0 changed, 0 destroyed.
```

```
Outputs:
```

```
ec2_instacne_private_ip = "192.168.0.5"
ec2_instacne_public_ip = "54.86.192.214"
```

Lecture 60

We should get tf output ec2_security_groups in a non redacted value from terraform.tfstate file.

```
ec2_instacne_private_ip = "192.168.0.11"
ec2_instacne_public_ip = "44.201.209.55"
ec2_security_groups = <sensitive>
```

```
BJIT@11728-Mahmuda-Akter-Keya MINGW64 /d/terraform tutorial
$ tf output ec2_security_groups
toset([])
```

Terraform output in json format:

```
$ tf output -json
{
  "ec2_instance_private_ip": {
    "sensitive": false,
    "type": "string",
    "value": "192.168.0.11"
  },
  "ec2_instance_public_ip": {
    "sensitive": false,
    "type": "string",
    "value": "44.201.209.55"
  },
  "ec2_security_groups": {
    "sensitive": true,
    "type": [
      "set",
      "string"
    ],
    "value": []
  }
}
```

Lecture 61

Terraform Local Values

Dry principal

Using dry principal, we can reduce the repetition of the code or we can reuse the code.

A local value assigns a **name to an expression**, so you can use that **name** multiple times within a module without repeating it.

Local values are like a **function's temporary local variables**.

Once a local value is declared, you can reference it in expressions as **local.<NAME>**.

```
locals {  
    service_name = "forum"  
    owner        = "Community Team"  
}  
  
locals {  
    # Common tags to be assigned to all resources  
    common_tags = {  
        Service = local.service_name  
        Owner   = local.owner  
    }  
}  
  
resource "aws_instance" "example" {  
    # ...  
  
    tags = local.common_tags  
}
```

Local values can be helpful to **avoid repeating the same values or expressions multiple times** in a configuration

If **overused** they can also make a configuration **hard to read** by future maintainers **by hiding the actual values used**

The ability to easily change the value in a central place is the **key advantage** of local values.

In short, Use local values only in moderation

Lecture 62

```
# local variable example
```

```
resource "aws_s3_bucket" "tf_s3" {  
  
    bucket = local.bucket-name  
  
    tags = {  
        Name      = local.bucket-name  
        Environment = var.environment_name  
    }  
}
```

```
#adding controller_access
```

```

resource "aws_s3_bucket_ownership_controls" "example" {

  bucket = local.bucket-name      #Required
  rule {
    object_ownership = "BucketOwnerPreferred" #Required
  }
  depends_on = [aws_s3_bucket.tf_s3]
}

resource "aws_s3_bucket_acl" "tf_s3_acl" {
  depends_on = [aws_s3_bucket_ownership_controls.example]

  bucket = local.bucket-name
  acl   = "private"
}

locals {
  bucket-name = "${var.app_name}-${var.environment_name}-bucket"
}

```

Lecture 63

Terarform Data Sources

Terraform Datasources

Data sources allow data to be fetched or computed for use elsewhere in Terraform configuration.

Use of data sources allows a Terraform configuration to make use of information defined outside of Terraform, or defined by another separate Terraform configuration.

A data source is accessed via a special kind of resource known as a *data resource*, declared using a data block

Terraform Datasources

We can refer the data resource in a resource as depicted

Meta-Arguments for Datasources

```
# Create EC2 Instance - Amazon Linux
resource "aws_instance" "my-ec2-vm" {
  ami           = data.aws_ami.amzlinux.id
  instance_type = var.ec2_instance_type
  key_name      = "terraform-key"
  user_data     = file("apache-install.sh")
  vpc_security_group_ids = [aws_security_group.id]
  tags = {
    "Name" = "amz-linux-vm"
  }
}
```

Data resources support the **provider** meta-argument as defined for managed resources, with the **same syntax and behavior**.

Data resources **do not currently have** any customization settings available for their **lifecycle**, but the **lifecycle nested block is reserved** in case any are added in future versions.

Data resources support **count** and **for_each** meta-arguments as defined for managed resources, with the **same syntax and behavior**.

Each instance will **separately read** from its data source with its own variant of the constraint arguments, producing an **indexed result**.

Datasource is Data block.

Each data resource is associated with a **single data source**, which determines the **kind of object (or objects)** it reads and what **query constraint arguments** are available

Data resources have the **same dependency resolution behavior** as defined for managed resources. Setting the **depends_on** meta-argument within data blocks **defers** reading of the data source until after all changes to the dependencies have been applied.

Reference:

- https://www.youtube.com/playlist?list=PLxBANRJzAw7i_xMJH-w2vzV8nP_AcIP6
- <https://developer.hashicorp.com/terraform/language/functions/file>
- <https://registry.terraform.io/providers/figma/aws-4-49-0/latest/docs/data-sources/ami>
- <https://registry.terraform.io/providers/figma/aws-4-49-0/latest/docs>
- <https://developer.hashicorp.com/terraform/language/functions/file>
- <https://developer.hashicorp.com/terraform/cli/commands/console>
- <https://spacelift.io/blog/terraform-console>
- <https://registry.terraform.io/providers/terraform-redhat/rhcs/latest/docs/guides/terraform-vars>
- <https://github.com/stacksimplify/hashicorp-certified-terraform-associate/blob/main/05-Terraform-Variables/05-01-Terraform-Input-Variables/README.md>
- <https://spacelift.io/blog/terraform-tfvars>
- <https://developer.hashicorp.com/terraform/language/values/variables>
- <https://developer.hashicorp.com/terraform/language/providers/configuration>
- <https://developer.hashicorp.com/terraform/language/providers/configuration>
- https://registry.terraform.io/providers/hashicorp/aws/latest/docs/resources/s3_bucket_acl
- https://registry.terraform.io/providers/BigEyeLabs/aws-test/latest/docs/resources/s3_bucket_ownership_controls
- <https://registry.terraform.io/providers/hashicorp/aws/latest/docs/resources/eip>
- https://developer.hashicorp.com/terraform/language/meta-arguments/for_each
- https://registry.terraform.io/providers/hashicorp/aws/latest/docs/resources/s3_bucket
- https://registry.terraform.io/providers/hashicorp/aws/latest/docs/resources/route_table
- <https://developer.hashicorp.com/terraform/language/expressions>
- <https://developer.hashicorp.com/terraform/language/meta-arguments/count>
- <https://developer.hashicorp.com/terraform/language/attr-as-blocks#arbitrary-expressions-with-argument-syntax>
- <https://developer.hashicorp.com/terraform/language/attr-as-blocks>
- <https://registry.terraform.io/providers/hashicorp/aws/latest/docs/resources/instance>
- <https://developer.hashicorp.com/terraform/tutorials/configuration-language/resource>
- <https://developer.hashicorp.com/terraform/tutorials/aws-get-started/aws-build>
- <https://developer.hashicorp.com/terraform/cli/commands/init#upgrade>
- <https://registry.terraform.io/providers/hashicorp/aws/3.74.3/docs/resources/vpc>

- <https://www.digitalocean.com/community/tutorials/how-to-create-reusable-infrastructure-with-terraform-modules-and-templates>
- <https://developer.hashicorp.com/terraform/language/resources/provisioners/syntax>